

# PREPARING DATA

---

Introducción a la Ciencia de Datos

# EDA steps

## Stage 1

- Frame the problem
- Ask interesting questions. What would you like to model?

## Stage 2

- Acquire the data, resources?
- Which are the relevant data?

## Stage 3

- Explore the data
- Plot the data, find trends and correlation patterns

## Stage 4

- Manipulate the data, solve problems missing values, outliers
- Imputation and transformation of data
- Feature engineering

# Prepare the ingredients: your data

- Reading data from CSV files
- Reading XML data
- Reading JSON data
- Reading data from fixed-width formatted files
- Reading data from R files and R libraries
- Removing cases with missing values
- Replacing missing values with the mean
- Removing duplicate cases
- Rescaling a variable to specified min-max range
- Normalizing or standardizing data in a data frame
- Binning numerical data
- Creating dummies for categorical variables
- Handling missing data
- Correcting data
- Imputing data
- Detecting outliers

# Handling missing data

- In most real-world problems, data is likely to be incomplete because of:
  - incorrect data entry,
  - faulty equipment,
  - improperly coded data.
- In R, missing values are represented by the symbol NA (not available) and are considered to be the first obstacle in predictive modeling.
- So, it's always a good idea to check for missing data in a dataset before proceeding for further predictive analysis.

# Why my data has missing values?

- **Data Extraction:** It is possible that there are problems with extraction process. In such cases, we should double-check for correct data
- **Data collection:** These errors occur at time of data collection and are harder to correct:
  - **Missing completely at random:** probability of missing variable is same for all observations.
  - **Missing at random:** missing ratio varies for different values / level of other input variables. (age data collection)
  - **Missing that depends on unobserved predictors:** missing values are not random and are related to the unobserved input variable. (discomfort)
  - **Missing that depends on the missing value itself:** This is a case when the probability of missing value is directly correlated with missing value itself (higher & lower income)

# Handling missing data

- R provides three simple ways to handle missing values:
  1. Deleting the observations.
  2. Deleting the variables.
  3. Transforming and binning values
  4. Replacing the values with mean, median, or mode.
  5. Prediction Model \*\*
  6. KNN Imputation \*\*

# Removing cases with missing values

- To get a data frame with no missing values for any variable, use the `na.omit()` function:
  - `dat.cleaned <- na.omit(dat)`
- We might sometimes want to selectively eliminate cases that have NA only for a specific variable. The example data frame has two missing values for Income. To get a data frame with only these two cases removed, use:

```
> dat.income.cleaned <- dat[!is.na(dat$Income),]  
> nrow(dat.income.cleaned)  
[1] 25
```

# Finding cases with no missing values

- The `complete.cases()` function takes a data frame or table as its argument and returns a Boolean vector with TRUE for rows that have no missing values, and FALSE otherwise:

```
> complete.cases(dat)
```

```
[1]  TRUE  TRUE  TRUE FALSE  TRUE FALSE  TRUE  TRUE  TRUE  
[10]  TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE FALSE  TRUE  
[19]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
```

Rows 4, 6, 13, and 17 have at least one missing value.

Instead of using the `na.omit()` function, we can do the following as well:

```
> dat.cleaned <- dat[complete.cases(dat),]  
> nrow(dat.cleaned)  
[1] 23
```



# Converting specific values to NA

- Sometimes, a specific value in a data frame means that the data was not available.
- For example, in the `dat` data frame, a value of 0 for `Income` probably means that the data is missing. We can convert these to `NA` by a simple assignment

```
> dat$Income[dat$Income==0] <- NA
```

# Handling missing data

## 4. Replacing the values with mean, median, or mode

- **Generalized Imputation:** we calculate the mean or median for all non missing values of that variable then replace missing value with mean or median.
- **Similar case Imputation:** In this case, we calculate average for gender “**Male**” (29.75) and “**Female**” (25) individually of non missing values then replace the missing value based on gender

## Replacing missing values with the mean

- Disregarding cases with any missing variables implies losing useful information.
- You may sometimes want to impute reasonable values (those that will not skew the results of analysis very much) for the missing values.
- replace missing values with the mean

```
> dat <- read.csv("missing-data.csv", na.strings = "")  
> dat$Income.imp.mean <- ifelse(is.na(dat$Income),  
mean(dat$Income, na.rm=TRUE), dat$Income)
```

# Understanding missing data

- use the `md.pattern()` function from the `mice` package to get a better understanding of the pattern of missing data.


```
> library(mice)
> install.packages("housingData")
> housing.dat<-housingData::housing
> str(housing.dat)
> md.pattern(housing.dat)
```

```
fips county state time medListPriceSqft medSoldPriceSqft nSold
67740      1      1      1      1              1          1      0
33453      1      1      1      1              1          1      1
11403      1      1      1      1              0          1      1
15929      1      1      1      1              1          1      1
 3259      1      1      1      1              0          1      2
109883     1      1      1      1              1          0      2
 5415      1      1      1      1              0          0      2
      0      0      0      0              20077      131227 146595 297899
```

# The package VIM

## (visualization and imputation of missing values)

```
library(VIM)
aggr_plot <- aggr(housing.dat, col=c('blue','red'), numbers=TRUE,
  sortVars=TRUE, labels=names(housing.dat), cex.axis=.7, gap=3,
  ylab=c("Histogram of missing data","Pattern"))
```



Next we will visualize the housing data to understand missing information

Variables sorted by number of missings:

Variable	Count
nSold	0.59330506
medSoldPriceSqft	0.53110708
medListPriceSqft	0.08125642
fips	0.00000000
county	0.00000000
state	0.00000000
time	0.00000000

# Identifying duplicate cases

- Sometimes we just want to identify the duplicated values without necessarily removing them for this, use the `duplicated()` function:

```
> duplicated(prospect)
```

```
[1] FALSE FALSE FALSE FALSE TRUE FALSE TRUE
```

- From the data, we know that cases 2, 5, and 7 are duplicates. To list the duplicate cases, use the following code:

```
> prospect[duplicated(prospect), ]
```

	salary	family.size	car
5	30000	3	Compact
7	30000	3	Compact

# Removing duplicate cases

- Removing duplicates

```
> salary <- c(20000, 30000, 25000, 40000, 30000, 34000, 30000)
> family.size <- c(4,3,2,2,3,4,3)
> car <- c("Luxury", "Compact", "Midsize", "Luxury", "Compact",
"Compact", "Compact")
> prospect <- data.frame(salary, family.size, car)
```

- The `unique()` function can do the job. It takes a vector or data frame as an argument and returns an object of the same type as its argument, but with duplicates removed.

```
> prospect
  salary family.size   car
1  20000           4 Luxury
2  30000           3 Compact
3  25000           2 Midsize
4  40000           2  Luxury
5  30000           3 Compact
6  34000           4 Compact
7  30000           3 Compact
```

```
> prospect.cleaned <- unique(prospect)
> nrow(prospect)
[1] 7
> nrow(prospect.cleaned)
[1] 5
```

# Rescaling a variable to specified min-max range

- Distance computations play a big role in many data analytics techniques.
- We know that variables with higher values tend to dominate distance computations and you may want to rescale the values between  $[0,1]$

```
> install.packages("scales")
➤ library(scales)
➤ library(mice)
> head(housing.dat)
> housing.dat$nSold.rescaled <- rescale(housing.dat$nSold)
```

fips	county	state	time	nSold	medListPriceSqft	medSoldPriceSqft	nSold.rescaled
1 06037	Los Angeles County	CA	2008-01-31	505900	NA	360.1645	0.3162685
2 06037	Los Angeles County	CA	2008-02-29	497100	NA	353.9788	0.3103794
3 06037	Los Angeles County	CA	2008-03-31	487300	NA	349.7633	0.3038212



# Rescaling a variable to specified min-max range

- To rescale use the function `rescale()` from the library “scales”

```
housing.dat
```

```
# rescale between [0,1]
```

```
housing.dat$nSold.rescaled <- rescale(housing.dat$nSold)
```

```
head(housing.dat)
```

```
# rescale to other ranges
```

```
housing.dat$nSold.rescaled
```

```
<- rescale(housing.dat$nSold, to = c(1, 100))
```

# Rescaling many variables at once

- Here an example of function to rescale many variables:

```
rescale.many <- function(dat, column.nos) {  
  nms <- names(dat)  
  for(col in column.nos) {  
    name <- paste(nms[col], ".rescaled", sep = "")  
    dat[name] <- rescale(dat[,col])  
  }  
  cat(paste("Rescaled ", length(column.nos), " variable(s)\n"))  
  dat  
}
```

- With the preceding function it is possible to rescale the fifth and seventh variables in the data frame:

```
> rescale.many(housing.dat, c(5,7))
```

# Normalizing or standardizing data in a data frame

- The `scale()` function computes the standard z score for each value (ignoring NAs) of each variable. That is, from each value it subtracts the mean and divides the result by the standard deviation of the associated variable.
- The `scale()` function takes two optional arguments, `center` and `scale`:

Argument	Effect
<code>center = TRUE, scale = TRUE</code>	Default behavior described earlier
<code>center = TRUE, scale = FALSE</code>	From each value, subtract the mean of the concerned variable
<code>center = FALSE, scale = TRUE</code>	Divide each value by the root mean square of the associated variable, where root mean square is $\sqrt{\text{sum}(x^2)/(n-1)}$
<code>center = FALSE, scale = FALSE</code>	Return the original values unchanged

# Binning numerical data

- Sometimes, we need to convert numerical data to categorical data or a factor:
  - *(e.g. Naive Bayes classification requires all variables to be categorical).*
  - In other situations, we may want to apply a classification method to a problem where the dependent variable is numeric but needs to be categorical

# Binning numerical data

Income is a numeric variable, and you may want to create a categorical variable from it by creating bins.

- Suppose you want to label incomes of \$10,000 or below as Low, incomes between \$10,000 and \$31,000 as Medium, and the rest as High

```
#Create a vector of break points:
> b <- c(-Inf, 10000, 31000, Inf)
#Create a vector of names for break points:
> names <- c("Low", "Medium", "High")
#Cut the vector using the break points with the function cut():
> students$Income.cat <- cut(students$Income, breaks = b, labels = names)
> head(students)
```

	Age	State	Gender	Height	Income	Income.cat
1	23	NJ	F	61	5000	Low
2	13	NY	M	55	1000	Low
3	58	NY	F	70	30000	Medium
4	29	TX	F	63	10000	Low

# Binning numerical data

If we leave out names, `cut()` uses the numbers in the second argument to construct interval names

```
#Create a vector of break points:
> b <- c(-Inf, 10000, 31000, Inf)
#Create a vector of names for break points:
> names <- c("Low", "Medium", "High")
#Cut the vector using the break points:
> students$Income.cat <- cut(students$Income, breaks = b)
> head(students)
```

	Age	State	Gender	Height	Income	Income.cat	Income.cat1
1	23	NJ	F	61	5000	Low	(-Inf,1e+04]
2	13	NY	M	55	1000	Low	(-Inf,1e+04]
3	36	NJ	M	66	3000	Low	(-Inf,1e+04]
4	31	VA	F	64	4000	Low	(-Inf,1e+04]
5	58	NY	F	70	30000	Medium	(1e+04,3.1e+04]
6	29	TX	F	63	10000	Low	(-Inf,1e+04]
7	39	NJ	M	67	50000	High	(3.1e+04, Inf]

## Creating a specified number of intervals automatically

- Rather than determining the breaks and hence the intervals manually, as mentioned earlier, we can specify the number of bins we want, say  $n$ , and let the `cut()` function handle the rest automatically.
- In this case, `cut()` creates  $n$  intervals of approximately equal width, as follows:

```
>students$Income.cat2 <- cut(students$Income,  
breaks = 4, labels = c("Level1", "Level2", "Level3","Level4"))
```

# Creating dummies for categorical variables

- This is needed in situations where we have categorical variables (factors) but need to use them in analytical methods that require numbers:
  - *K nearest neighbors (KNN)*
  - *Linear Regression*

```
> install.packages("dummies")  
> library(dummies)  
> attach(Salaries)  
> head(Salaries)
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
1		Prof	B	19	18	Male 139750
2		Prof	B	20	16	Male 173200
3		AsstProf	B	4	3	Male 79750



# Creating dummies for categorical variables

```
> names(Salaries)
[1] "rank"          "discipline"    "yrs.since.phd" "yrs.service"
"sex"  "salary"
str(Salaries)
data.frame':   397 obs. of  6 variables:
 $ rank          : Factor w/ 3 levels "AsstProf","AssocProf",...: 3 3  $
discipline      : Factor w/ 2 levels "A","B": 2 2 2 2 2 2 2 2 2 2
 $ yrs.since.phd: int   19 20 4 45 40 6 30 45 21 18 ...
 $ yrs.service  : int   18 16 3 39 41 6 23 45 20 18 ...
 $ sex          : Factor w/ 2 levels "Female","Male": 2 2 2 2 2 2 2 2 $
salary         : int  139750 173200 79750 115000 141500 97000
```

```
Salaries.dummy<- dummy.data.frame(Salaries, sep = ".")
```

# Creating dummies for categorical variables

```
names(Salaries.dummy)
```

```
[1] "rank.AsstProf" "rank.AssocProf" "rank.Prof"      "discipline.A"
"discipline.B"  "yrs.since.phd"  "yrs.service"    "sex.Female"
"sex.Male"      "salary"
```

```
> head(Salaries.dummy)
```

	rank.AsstProf	rank.AssocProf	rank.Prof	discipline.A	discipline.B	yrs.since.phd	yrs.service
1	0	0	1	0	1	19	18
2	0	0	1	0	1	20	16
3	1	0	0	0	1	4	3
4	0	0	1	0	1	45	39
5	0	0	1	0	1	40	41
6	0	1	0	0	1	6	6

	sex.Female	sex.Male	salary
1	0	1	139750
2	0	1	173200
3	0	1	79750
4	0	1	115000
5	0	1	141500
6	0	1	97000

# Detecting outliers

Outlier is an observation that appears far away and diverges from an overall pattern in a sample. Outlier can be of two types:

- Univariate
- Multivariate.

The ideal way to deal with them is to find out the reason of having these outliers. Causes of outliers can be classified in two broad categories:

- Artificial (Error) / Non-natural
- Natural.

# Cleaning data

The five most common problems with messy datasets

- Column headers are values, not variable names.
- Multiple variables are stored in one column.
- Variables are stored in both rows and columns.
- Multiple types of observational units are stored in the same table.
- A single observational unit is stored in multiple tables.

# Cleanning data packages in R

- **dplyr**: package for data manipulation
- **tidyr**: suitable for data tidying (splitting, joining variables) and it supports the dplyr package's data pipelines
- **reshape2**: it provides easy aggregations and restructuring of data
- **data.table**: allows fast aggregations and transformations of large data sets

Some packages specialize in addressing character variables, strings, and text or date and time information are:

- **stringi and stringr** support work with strings and allow, for example, pattern searching, string generation, date and time formatting, parsing, ect...
- **lubridate**: helps R users with date and time stamps, their extraction and advanced manipulation.

## tidyr()

- `gather()` and `spread()` convert data between wide and long format
- `separate()` and `unite()` separate a single column into multiple columns and vice versa
- `complete()` turns implicit missing values in explicit missing values by completing missing data combinations

# gather()

```
> messy
```

	id	trt	work.T1	home.T1	work.T2	home.T2
1	1	treatment	0.08513597	0.6158293	0.1135090	0.05190332
2	2	control	0.22543662	0.4296715	0.5959253	0.26417767
3	3	treatment	0.27453052	0.6516557	0.3580500	0.39879073
4	4	control	0.27230507	0.5677378	0.4288094	0.83613414

We want to gather all the columns, except for the id and trt ones, in two columns key and value

```
gathered.messy <-  
gather(messy, key, value, -id, -trt)
```

	id	trt	key	value
1	1	treatment	work.T1	0.08513597
2	2	control	work.T1	0.22543662
3	3	treatment	work.T1	0.27453052
4	4	control	work.T1	0.27230507
5	1	treatment	home.T1	0.61582931
6	2	control	home.T1	0.42967153

# spread()

Takes different levels of a factor and spreads them out into different columns. This means we can convert from long data to wide.

```
> stocksm
```

	time	stock	price
1	2009-01-01	X	-0.66184983
2	2009-01-02	X	1.71895416
3	2009-01-03	X	2.12166699
4	2009-01-04	Y	1.49715368
5	2009-01-05	Y	-0.03614058
6	2009-01-06	Y	1.23194518

```
spread.stock <- spread(stocksm, stock, price)
```

```
head(spread.stock)
```

	time	X	Y	Z
1	2009-01-01	-0.66184983	-0.7656438	-5.0672590
2	2009-01-02	1.71895416	0.5988432	-0.7943331
3	2009-01-03	2.12166699	1.3484795	0.5554631



# Piping and chaining code

- dplyr allows you to use the pipe (`%>%`) operator to chain functions together.
- Chaining code allows you to streamline your workflow and make it easier to read.
- When using the `%>%` operator, first specify the data frame that all following functions will use.
- For the rest of the chain the data frame argument can be omitted from the remaining functions.

# unite()

- `unite(data, col, ..., sep = "_", remove = TRUE)`

```
> data
```

	date	hour	min	second	event
1	2016-01-01	7	30	29	u
2	2016-01-02	9	43	36	a
3	2016-01-03	13	58	60	l
4	2016-01-04	20	22	11	q
5	2016-01-05	5	44	47	p

```
> dataNew
```

	datetime	event
1	2016-01-01 7:30:29	u
2	2016-01-02 9:43:36	a
3	2016-01-03 13:58:60	l
4	2016-01-04 20:22:11	q

```
# Now, let us combine the date, hour, min, and  
# second columns into a new column called datetime.
```

```
dataNew <- data %>%  
  unite(datehour, date, hour, sep = ' ') %>%  
  unite(datetime, datehour, min, second, sep = ':')  
dataNew
```

## separate()

- We can get back the original data we created using `separate` as follows:

```
data1 <- dataNew %>%  
  separate(datetime, c('date', 'time'), sep = ' ') %>%  
  separate(time, c('hour', 'min', 'second'), sep = ':')  
data1
```

# dyplyr()

- `filter()` subset data based on logical criteria
- `select()` select certain columns
- `arrange()` order rows by value of a column
- `rename()` rename columns
- `group_by()` group data by common variables for performing calculations
- `mutate()` create a new variable/column
- `summarize()` summarize data into a single row of values

## filter()

- The filter function will return all the rows that satisfy a following condition.

```
filter(airquality, Temp > 70)
```

```
filter(airquality, Temp > 80 & Month > 5)
```

# mutate()

- Mutate is used to add new variables to the data. For example, let's add a new column that displays the temperature in Celsius

```
mutate(airquality, TempInC = (Temp - 32) * 5 / 9)
```

# Summarise

- The summarise function is used to summarise multiple values into a single value. It is very powerful when used in conjunction with the other functions in the dplyr package, as demonstrated below. `na.rm = TRUE` will remove all NA values while calculating the mean, so that it doesn't produce spurious results.

```
> summarise(airquality, mean(Temp, na.rm = TRUE))  
  mean(Temp, na.rm = TRUE)  
1                77.88235
```

## Group by

- The `group_by` function is used to group data by one or more variables.
- For example, we can group the data together based on the Month, and then use the `summarise` function to calculate and display the mean temperature for each month.

```
summarise(group_by(airquality, Month), mean(Temp, na.rm = TRUE))  
# A tibble: 5 x 2  
  Month `mean(Temp, na.rm = TRUE)`  
  <int>           <dbl>  
1     5     65.54839  
2     6     79.10000  
3     7     83.90323  
4     8     83.96774  
5     9     76.90000
```



## arrange()

- The arrange function is used to arrange rows by variables. Currently, the airquality dataset is arranged based on Month, and then Day.
- We can arrange the rows in the descending order of Month, and then in the ascending order of Day.

```
arrange(airquality, desc(Month), Day)
```

# All together

```
msleep %>%  
  group_by(order) %>%  
  summarise(avg_sleep = mean(sleep_total),  
            min_sleep = min(sleep_total),  
            max_sleep = max(sleep_total),  
            total = n())  
  
## Source: local data frame [19 x 5]  
##  
##      order avg_sleep min_sleep max_sleep total  
## 1 Afrosoricida 15.600000    15.6    15.6      1  
## 2 Artiodactyla  4.516667     1.9     9.1      6  
## 3 Carnivora    10.116667     3.5    15.8     12  
## 4 Cetacea      4.500000     2.7     5.6      3  
## 5 Chiroptera   19.800000    19.7    19.9      2  
## 6 Cingulata    17.750000    17.4    18.1      2  
## 7 Didelphimorphia 18.700000    18.0    19.4      2  
## 8 Diprotodontia 12.400000    11.1    13.7      2
```